

Lab 3: Design-Based Statistical Inference for Causal Quantities

Dias Akhmetbekov

2024-02-16

Outline

- ▶ Inference for SATE and PATE
- ▶ Finite sample adjustments

Complete randomization

- ▶ Setting: **completely randomized experiment**
 - ▶ N units: N_1 treated and N_0 control
 - ▶ Random assignment: $D_i \perp (Y_{1i}, Y_{0i})$
- ▶ What we observe is $Y_i = D_i Y_{1i} + (1 - D_i) Y_{0i}$
- ▶ An intuitive estimator for the SATE:

$$\hat{\tau}_{\text{diff}} = \underbrace{\frac{1}{N_1} \sum_{i=1}^N D_i Y_i}_{\text{mean among treated}} - \underbrace{\frac{1}{N_0} \sum_{i=1}^N (1 - D_i) Y_i}_{\text{mean among control}}$$

- ▶ Conditional on the sample \mathcal{S} , $\hat{\tau}_{\text{diff}}$ only varies because of D_i

Complete randomization

- First, take expectations over repeated treatment assignments holding the sample \mathcal{S} fixed:

$$\begin{aligned} E_D[\hat{\tau}_{diff} \mid \mathcal{S}] &= \frac{1}{N_1} \sum_{i=1}^N [E_D[D_i Y_i \mid \mathcal{S}]] - \frac{1}{N_0} \sum_{i=1}^N [E_D[(1 - D_i) Y_i \mid \mathcal{S}]] \\ &= \frac{1}{N_1} \sum_{i=1}^N [E_D[D_i Y_i(1) \mid \mathcal{S}]] - \frac{1}{N_0} \sum_{i=1}^N [E_D[(1 - D_i) Y_i(0) \mid \mathcal{S}]] \\ &= \frac{1}{N_1} \sum_{i=1}^N [E_D[D_i \mid \mathcal{S}] Y_i(1)] - \frac{1}{N_0} \sum_{i=1}^N [E_D[(1 - D_i) \mid \mathcal{S}] Y_i(0)] \\ &= \frac{1}{N_1} \sum_{i=1}^N \left(\frac{N_1}{N} \right) Y_i(1) - \frac{1}{N_0} \sum_{i=1}^N \left(\frac{N_0}{N} \right) Y_i(0) \\ &= \frac{1}{N} \sum_{i=1}^N [Y_i(1) - Y_i(0)] = SATE \end{aligned}$$

- The unbiasedness results does NOT depend on N_0, N_1

Complete randomization

- ▶ Then take expectations over different samples:

$$\begin{aligned} E [E_D [\hat{\tau}_{diff} \mid \mathcal{S}]] &= E [SATE] \\ &= \frac{1}{N} \sum_{i \in \mathcal{S}} E [Y_{1i} - Y_{0i}] = \frac{1}{N} NE [Y_{1i} - Y_{0i}] \\ &= E [Y_{1i} - Y_{0i}] = PATE \end{aligned}$$

- ▶ $\hat{\tau}_{diff}$ is unbiased for PATE (under random sampling and random treatment assignment)

Finite-sample sampling variance

- Conditional variance for complete random assignment:

$$\begin{aligned}\text{Var}_D[\hat{\tau}_{diff} \mid \mathcal{S}] &= \text{Var}_D[\bar{Y}_1 \mid \mathcal{S}] + \text{Var}_D[\bar{Y}_0 \mid \mathcal{S}] - 2\text{Cov}_D[\bar{Y}_1, \bar{Y}_0 \mid \mathcal{S}] \\&= \frac{s_{Y_1}^2}{N_1} \left(\frac{N - N_1}{N} \right) + \frac{s_{Y_0}^2}{N_0} \left(\frac{N - N_0}{N} \right) - 2 \left[-\frac{s_{Y_1, Y_0}}{N} \right] \\&= \frac{s_{Y_1}^2}{N_1} + \frac{s_{Y_0}^2}{N_0} - \frac{s_{Y_1}^2 + s_{Y_0}^2 - 2s_{Y_1, Y_0}}{N} \\&= \frac{s_{Y_1}^2}{N_1} + \frac{s_{Y_0}^2}{N_0} - \frac{s_{\tau}^2}{N},\end{aligned}$$

- Last term is the in-sample variation of the individual treatment effects (not observable!)

Inferences from sampling variance

- ▶ Usual variance estimator is the Neyman (or robust) estimator:

$$\hat{V} = \frac{\hat{s}_0^2}{N_0} + \frac{\hat{s}_1^2}{N_1}$$

- ▶ \hat{V} is biased for $\text{Var}_D[\hat{\tau}_{diff} \mid \mathcal{S}]$
- ▶ It leads to conservative inferences:
 - ▶ standard errors $\sqrt{\hat{V}}$ will be at least as big as they should be;
 - ▶ CIs using $\sqrt{\hat{V}}$ will be at least wide as they should be;
 - ▶ Type I error rates will still be correct, power will be lower;
 - ▶ Both will be exactly right if treatment effects are constant.

Population variance

- From ANOVA theorem and CEF decomposition:
 $\text{Var}[A] = E[\text{Var}[A \mid B]] + \text{Var}[E[A \mid B]]$. Hence,

$$\begin{aligned}\text{Var}[\hat{\tau}] &= E[\text{Var}_D[\hat{\tau} \mid \mathcal{S}]] + \text{Var}[E_D[\hat{\tau} \mid \mathcal{S}]] \\&= E \left[\frac{s_{Y_1}^2}{N_1} + \frac{s_{Y_0}^2}{N_0} - \frac{s_{\tau}^2}{N} \right] + \text{Var}[SATE] \\&= E \left[\frac{s_{Y_1}^2}{N_1} + \frac{s_{Y_0}^2}{N_0} - \frac{s_{\tau}^2}{N} \right] + \text{Var} \left[\frac{1}{N} \sum_{i \in \mathcal{S}} \tau \right] \\&= E \left[\frac{s_{Y_1}^2}{N_1} + \frac{s_{Y_0}^2}{N_0} - \frac{s_{\tau}^2}{N} \right] + \frac{\sigma_{\tau}^2}{N} \\&= \frac{\sigma_{Y_1}^2}{N_1} + \frac{\sigma_{Y_0}^2}{N_0} - \frac{\sigma_{\tau}^2}{N} + \frac{\sigma_{\tau}^2}{N} \\&= \frac{\sigma_{Y_1}^2}{N_1} + \frac{\sigma_{Y_0}^2}{N_0}.\end{aligned}$$

Population variance

- ▶ The Neyman estimator \hat{V} is now unbiased for \hat{V}_{diff} ;
- ▶ Consistency via the law of large numbers: $\hat{\tau}_{\text{diff}} \xrightarrow{P} \text{PATE}$
- ▶ Asymptotic normality via the Central Limit Theorem (CLT):

$$\frac{\hat{\tau}_{\text{diff}} - \text{PATE}}{\sqrt{\frac{\sigma_1^2}{N_1} + \frac{\sigma_0^2}{N_0}}} \xrightarrow{d} \mathcal{N}(0, 1)$$

- ▶ Two interpretations for \hat{V} :
 - ▶ Unbiased estimator for sampling variance of the traditional estimator of the PATE;
 - ▶ Conservative estimator for the sampling variance of the traditional estimator of the SATE

Simulation

```
set.seed(123) # Set seed

# Function to generate population data
generate_population <- function(N_pop) {
  Y0 <- abs(rnorm(N_pop, mean = 5, sd = 2))
  Y1 <- Y0 + rnorm(N_pop, 0, 5) + 4
  TE <- Y1 - Y0
  data.frame(Y0 = Y0, Y1 = Y1, TE = TE)
}
```

Simulation

```
# Function to simulate the randomization process
simulate_randomization <- function(sample, Nboot) {
  Nsample <- nrow(sample)
  dim <- vars <- rep(NA, Nboot)
  for(i in 1:Nboot){
    D <- sample(c(0,1), Nsample, replace = TRUE)
    Y <- D*sample$Y1 + (1-D)*sample$Y0
    dim[i] <- mean(Y[D==1]) - mean(Y[D==0])
    vars[i] <- var(Y[D==1])/(Nsample/2) + var(Y[D==0])/(Nsample/2)
  }
  list(dim = dim, vars = vars)
}
```

Simulation

```
# Generate population
```

```
N_pop <- 100000
```

```
pop <- generate_population(N_pop)
```

```
(PATE <- mean(pop$TE))
```

```
## [1] 4.026077
```

```
# Single sample simulation
```

```
Nsample <- 1000
```

```
sample <- pop[sample(N_pop, Nsample),]
```

```
(SATE <- mean(sample$TE))
```

```
## [1] 3.910847
```

Simulation

```
# Randomization distribution simulation  
Nboot <- 1000  
randomization_results <- simulate_randomization(sample, Nboot)  
(mean_dim <- mean(randomization_results$dim))
```

```
## [1] 3.919121
```

```
(var_dim <- var(randomization_results$dim))
```

```
## [1] 0.03871511
```

```
(expected_vars <- mean(randomization_results$vars))
```

```
## [1] 0.06289298
```

Simulation

```
# Two-level simulation: sampling and randomization
Nsampling <- 100
dim_sampling <- matrix(NA, Nsampling, Nboot)
vars_sampling <- matrix(NA, Nsampling, Nboot)

for(j in 1:Nsampling){
  sample <- pop[sample(N_pop, Nsample),]
  results <- simulate_randomization(sample, Nboot)
  dim_sampling[j,] <- results$dim
  vars_sampling[j,] <- results$vars
}
```

Simulation

```
(mean_dim_sampling <- mean(dim_sampling))
```

```
## [1] 4.006696
```

```
(var_dim_sampling <- var(as.vector(dim_sampling)))
```

```
## [1] 0.06385131
```

```
(expected_vars_sampling <- mean(vars_sampling))
```

```
## [1] 0.06597705
```

```
# Theoretical variance
```

```
(theoretical_var <- var(pop$Y1)/(Nsample/2) + var(pop$Y0)/(Nsample/2))
```

```
## [1] 0.06595247
```

Finite samples

- ▶ Suppose the goal is robust inference for $\hat{\tau} = \bar{Y}_1 - \bar{Y}_0$ in a small sample size (30 units, for example)
- ▶ Suppose N_1 is very large but N_0 is very small
- ▶ Then, \bar{Y}_1 will be very precisely estimated, but \bar{Y}_0 will be imprecisely estimated
- ▶ That being the case, using $N - k = (N_1 + N_0) - k$ as the degrees of freedom adjustment overstates stability
- ▶ Correct degrees of freedom adjustment ought to be closer to $N_0 - k$

Simulation (Imbens and Kolesar, 2016)

```
set.seed(123)
library(sandwich)
library(lmtest)
library(MASS)
library(ggplot2)
```

```
# Simulation parameters
N <- 30
N1 <- 3
N0 <- N - N1
beta <- c(0, 0)
sigma_1 <- 1
sigma_0_values <- c(0.5, 0.85, 1, 1.18, 2)
num_replications <- 10000
true_effect <- beta[2]
```

Simulation (Imbens and Kolesar, 2016)

```
# Function to simulate data
simulate_data <- function(N, N1, sigma_0, sigma_1) {
  D <- c(rep(1, N1), rep(0, N - N1))
  Y <- rep(NA, N)
  epsilon <- rnorm(N, mean = 0, sd = ifelse(D == 1, sigma_1, sigma_0))
  Y <- beta[1] + beta[2] * D + epsilon
  data.frame(Y = Y, D = D)
}

# Function to calculate and evaluate CI coverage
calculate_coverage <- function(sigma_0, estimator_function) {
  coverage_count <- 0
  for (i in 1:num_replications) {
    data <- simulate_data(N, N1, sigma_0, sigma_1)
    model <- lm(Y ~ D, data = data)
    ci <- estimator_function(model)
    if (true_effect >= ci[1] && true_effect <= ci[2]) {
      coverage_count <- coverage_count + 1
    }
  }
  coverage_probability <- coverage_count/num_replications
  return(coverage_probability)
}
```

Simulation (Imbens and Kolesar, 2016)

```
# Estimator functions
```

```
homoskedastic_estimator <- function(model) {  
  ci <- confint(model, level = 0.95)["D",]  
  return(ci)  
}
```

```
heteroskedastic_estimator <- function(model) {  
  se <- sqrt(diag(vcovHC(model, type = "HC1")))  
  coef <- coef(model)["D"]  
  ci <- coef + qt(c(0.025, 0.975), df = N - 2) * se[2]  
  return(ci)  
}
```

Simulation (Imbens and Kolesar, 2016)

```
hc2_estimator <- function(model) {  
  se <- sqrt(diag(vcovHC(model, type = "HC2")))  
  coef <- coef(model)["D"]  
  ci <- coef + qt(c(0.025, 0.975), df = N - 2) * se[2]  
  return(ci)  
}  
  
bell_mccaffrey_estimator <-function(model) {  
  se <- sqrt(diag(vcovHC(model, type = "HC2")))  
  coef <- coef(model)["D"]  
  ci <- coef + qt(c(0.025, 0.975), df = N1 - 1) * se[2]  
  return(ci)  
}
```

Simulation (Imbens and Kolesar, 2016)

```
# Run simulations
coverage_results <- list()
for (sigma_0 in sigma_0_values) {
  coverage_homoskedastic <- calculate_coverage(sigma_0, homoskedastic_estimator)
  coverage_heteroskedastic <- calculate_coverage(sigma_0, heteroskedastic_estimator)
  coverage_HC2 <- calculate_coverage(sigma_0, hc2_estimator)
  coverage_BM <- calculate_coverage(sigma_0, bell_mccaffrey_estimator)
  coverage_results[[as.character(sigma_0)]] <- list(
    homo = coverage_homoskedastic,
    EHW = coverage_heteroskedastic,
    HC2 = coverage_HC2,
    BM = coverage_BM
  )
}
```

Simulation (Imbens and Kolesar, 2016)

```
# Initialize an empty data frame
coverage_df <- data.frame(Sigma_0 = character(),
                          Estimator = character(),
                          Coverage = numeric(),
                          stringsAsFactors = FALSE)

# Fill the data frame with coverage results
for (sigma_0 in names(coverage_results)) {
  for (estimator in names(coverage_results[[sigma_0]])) {
    new_row <- data.frame(Sigma_0 = sigma_0,
                          Estimator = estimator,
                          Coverage = coverage_results[[sigma_0]][[estimator]],
                          stringsAsFactors = FALSE)
    coverage_df <- rbind(coverage_df, new_row)
  }
}
```

Coverage Probabilities by Sigma_0 and Estimator

